

Project 3 - Report

Project Title: Deep Learning and Support Vector Machine Project

Objective:

To build two separate models:

1. **Classification Model:** A Convolutional Neural Network trained to classify mushroom images into one of 9 classes using a softmax output layer.
2. **Feature Extraction + SVM Model:** A CNN-based feature extractor that feeds its output into an SVM classifier for classification.

Dataset:

- The dataset consists of labeled mushroom images with 9 classes:

- Agaricus
- Amanita
- Boletus
- Cortinarius
- Entoloma
- Hygrocybe
- Lactarius
- Russula
- Suillus

- The training data is in the folder structure as below.

Mushrooms

- Agaricus
- Amanita
- Boletus
- Cortinarius
- Entoloma
- Hygrocybe
- Lactarius
- Russula
- Suillus

- The test data consists of images and a CSV file with image paths and labels as below.

image_path,label → mushrooms_test/test1.jpg,Amanita

Saved Model:

- Classification Model → classifier.keras
- Feature Extraction + SVM Model → svm_classifier.pkl

Model Explanation:**1) Model 1: CNN Classification Model (with Softmax)****Purpose:**

This model is used for end-to-end classification of mushroom images into one of 9 species.

Architecture Overview:

- Input: Images resized to 128×128×3
- Layers:
 - Data Augmentation (flip, rotate, zoom, brightness, etc.)
 - Convolutional Layers: 4 layers with increasing filters (32 → 64 → 128 → 128), using ReLU activation.
 - Pooling Layers: MaxPooling2D after each conv layer to reduce spatial dimensions.
 - GlobalAveragePooling2D to flatten before dense layers.
 - Dense Layer with 256 units + Dropout (0.5) for regularization.
 - Output Layer: Softmax activation with 9 units (for 9 classes).

Output: A probability distribution across 9 classes — the class with the highest probability is the predicted label.

Loss Function: categorical_crossentropy

Metric: accuracy

2) Model 2: Feature Extractor + SVM Classifier

Purpose:

To extract high-level features from images using the CNN model and classify them using an SVM (Support Vector Machine) instead of softmax.

Architecture Overview:

- Uses the same CNN layers as Model 1 up to the last dropout layer, but removes the softmax classification layer.
- Output of the model is a 256-dimensional feature vector (from the last dense layer).

Pipeline:

1. Pass images through the CNN feature extractor.
2. Collect the feature vectors (`X_train`) and corresponding labels (`y_train`).
3. Train a Support Vector Machine (SVM) classifier (`SVC(kernel='linear')`) on these features.
4. Evaluate SVM on extracted validation features

Approach:

1. Load and preprocess image data

- Used `tf.keras.utils.image_dataset_from_directory` to load images from the dataset folder.
- Applied resizing to **128x128** resolution and used categorical labels for multi-class classification.
- Split the dataset into 80% training and 20% validation using the `validation_split` argument.

2. Apply data augmentation

- Created a custom `Sequential` pipeline using `tf.keras.layers` for transformations such as:
 - Horizontal flipping
 - Rotation
 - Zoom
 - Brightness and contrast adjustment
 - Translation
- This helps the model generalize better and prevent overfitting.



3. Build a custom CNN model with softmax output

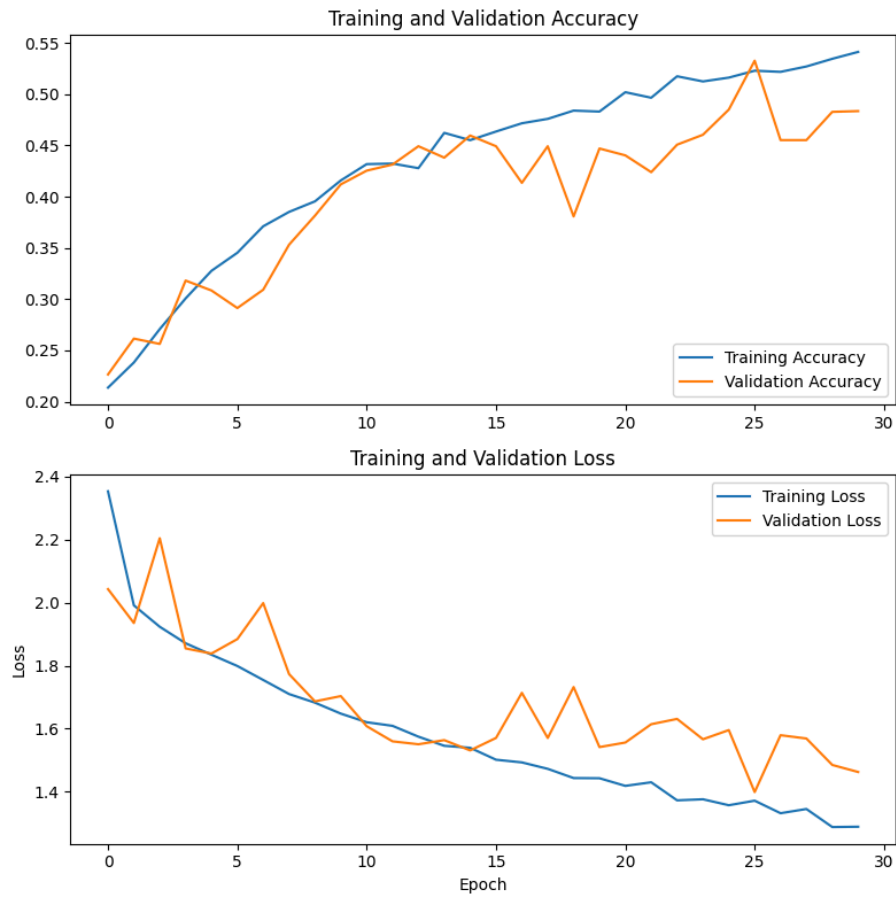
- Constructed a CNN from scratch using `Conv2D`, `MaxPooling2D`, `GlobalAveragePooling2D`, and `Dense` layers.
- Final layer is a `Dense` layer with `softmax` activation for 9-class classification.
- Included `Dropout` (0.5) before the output layer to reduce overfitting.

| Layer (type) | Output Shape | Param # |
|--------------------------------------------------------------------|----------------------|---------|
| input_layer_16 (<code>InputLayer</code>) | (None, 128, 128, 3) | 0 |
| sequential_7 (<code>Sequential</code>) | (None, 128, 128, 3) | 0 |
| conv2d_32 (<code>Conv2D</code>) | (None, 128, 128, 32) | 896 |
| max_pooling2d_24 (<code>MaxPooling2D</code>) | (None, 64, 64, 32) | 0 |
| conv2d_33 (<code>Conv2D</code>) | (None, 64, 64, 64) | 18,496 |
| max_pooling2d_25 (<code>MaxPooling2D</code>) | (None, 32, 32, 64) | 0 |
| conv2d_34 (<code>Conv2D</code>) | (None, 32, 32, 128) | 73,856 |
| max_pooling2d_26 (<code>MaxPooling2D</code>) | (None, 16, 16, 128) | 0 |
| conv2d_35 (<code>Conv2D</code>) | (None, 16, 16, 128) | 147,584 |
| global_average_pooling2d_8 (<code>GlobalAveragePooling2D</code>) | (None, 128) | 0 |
| dense_16 (<code>Dense</code>) | (None, 256) | 33,024 |
| dropout_8 (<code>Dropout</code>) | (None, 256) | 0 |
| dense_17 (<code>Dense</code>) | (None, 9) | 2,313 |

4. Train the CNN model

- Compiled with **ADAM** optimizer and **categorical_crossentropy** loss function.
- Trained the model for 30 epochs using `model.fit()` with augmented training and validation datasets.
- Saved the model in `.keras` format for reuse: `model.save("trial_model.keras")`.

```
Epoch 1/30
167/167 ----- 116s 622ms/step - accuracy: 0.1980 - loss: 3.9597 - val_accuracy: 0.2303 - val_loss: 2.0340
Epoch 2/30
167/167 ----- 118s 701ms/step - accuracy: 0.2315 - loss: 2.0061 - val_accuracy: 0.2474 - val_loss: 1.9620
Epoch 3/30
167/167 ----- 108s 641ms/step - accuracy: 0.2640 - loss: 1.9454 - val_accuracy: 0.2861 - val_loss: 1.9385
Epoch 4/30
167/167 ----- 98s 585ms/step - accuracy: 0.3039 - loss: 1.8854 - val_accuracy: 0.3212 - val_loss: 1.8805
Epoch 5/30
167/167 ----- 109s 653ms/step - accuracy: 0.3148 - loss: 1.8664 - val_accuracy: 0.2988 - val_loss: 1.8839
Epoch 6/30
167/167 ----- 110s 654ms/step - accuracy: 0.3137 - loss: 1.8471 - val_accuracy: 0.2705 - val_loss: 1.8973
Epoch 7/30
167/167 ----- 100s 598ms/step - accuracy: 0.3434 - loss: 1.8014 - val_accuracy: 0.3465 - val_loss: 1.7638
Epoch 8/30
167/167 ----- 95s 570ms/step - accuracy: 0.3758 - loss: 1.7350 - val_accuracy: 0.3025 - val_loss: 1.8914
Epoch 9/30
167/167 ----- 101s 602ms/step - accuracy: 0.3766 - loss: 1.7111 - val_accuracy: 0.3696 - val_loss: 1.7240
Epoch 10/30
167/167 ----- 83s 498ms/step - accuracy: 0.3975 - loss: 1.6747 - val_accuracy: 0.3793 - val_loss: 1.7704
Epoch 11/30
167/167 ----- 76s 452ms/step - accuracy: 0.4014 - loss: 1.6516 - val_accuracy: 0.4367 - val_loss: 1.5947
Epoch 12/30
167/167 ----- 76s 456ms/step - accuracy: 0.4205 - loss: 1.6216 - val_accuracy: 0.4143 - val_loss: 1.6578
Epoch 13/30
...
Epoch 29/30
167/167 ----- 65s 385ms/step - accuracy: 0.5486 - loss: 1.2779 - val_accuracy: 0.5171 - val_loss: 1.4795
Epoch 30/30
167/167 ----- 67s 400ms/step - accuracy: 0.5495 - loss: 1.2578 - val_accuracy: 0.4747 - val_loss: 1.5388
```



5. Build a CNN-based feature extractor model

- Removed the final softmax layer and used the CNN base up to the dense representation layer (256 units).
- This model outputs feature vectors instead of class predictions.

| Layer (type) | Output Shape | Param # |
|------------------------------------------------------|----------------------|---------|
| input_layer_18 (InputLayer) | (None, 128, 128, 3) | 0 |
| sequential_7 (Sequential) | (None, 128, 128, 3) | 0 |
| conv2d_40 (Conv2D) | (None, 128, 128, 32) | 896 |
| max_pooling2d_30 (MaxPooling2D) | (None, 64, 64, 32) | 0 |
| conv2d_41 (Conv2D) | (None, 64, 64, 64) | 18,496 |
| max_pooling2d_31 (MaxPooling2D) | (None, 32, 32, 64) | 0 |
| conv2d_42 (Conv2D) | (None, 32, 32, 128) | 73,856 |
| max_pooling2d_32 (MaxPooling2D) | (None, 16, 16, 128) | 0 |
| conv2d_43 (Conv2D) | (None, 16, 16, 128) | 147,584 |
| global_average_pooling2d_10 (GlobalAveragePooling2D) | (None, 128) | 0 |
| dense_19 (Dense) | (None, 256) | 33,024 |
| dropout_10 (Dropout) | (None, 256) | 0 |

6. Extract features and train an SVM

- Used the CNN feature extractor to generate fixed-size feature vectors for both training and validation datasets.
- Flattened labels from one-hot to class indices using `np.argmax()`.

```
1/1 ————— 0s 127ms/step
1/1 ————— 0s 68ms/step
1/1 ————— 0s 77ms/step
1/1 ————— 0s 72ms/step
1/1 ————— 0s 69ms/step
1/1 ————— 0s 75ms/step
1/1 ————— 0s 65ms/step
1/1 ————— 0s 70ms/step
1/1 ————— 0s 75ms/step
1/1 ————— 0s 70ms/step
1/1 ————— 0s 76ms/step
1/1 ————— 0s 80ms/step
1/1 ————— 0s 78ms/step
1/1 ————— 0s 78ms/step
1/1 ————— 0s 74ms/step
1/1 ————— 0s 62ms/step
1/1 ————— 0s 70ms/step
1/1 ————— 0s 75ms/step
1/1 ————— 0s 70ms/step
1/1 ————— 0s 68ms/step
1/1 ————— 0s 66ms/step
1/1 ————— 0s 99ms/step
1/1 ————— 0s 69ms/step
1/1 ————— 0s 72ms/step
1/1 ————— 0s 80ms/step
...
1/1 ————— 0s 92ms/step
1/1 ————— 0s 101ms/step
1/1 ————— 0s 100ms/step
1/1 ————— 0s 104ms/step
```

7. Train and evaluate an SVM classifier

- Used [scikit-learn's SVC](#) with a linear kernel.

- Trained it on the CNN-generated features.
- Evaluated the SVM model on validation features using `accuracy_score`.
- Saved the trained model as `svm_classifier.pkl` using `joblib`.

8. Visualize training performance

- Plotted training and validation accuracy and loss over epochs using `matplotlib`.

```
Fitting 3 folds for each of 15 candidates, totalling 45 fits  
SVM classifier saved as svm_classifier.pkl  
Best parameters: {'C': 1, 'max_iter': 1000}
```

```
SVM Classifier Accuracy: 0.4396
```

Testing:

The `proj3_classification_test.py` and `proj3_extractSVM_test.py` script loads the saved model and processes the test CSV and images.

Requirements:

```
python==3.11.0  
pip==22.3  
tensorflow==2.19.0  
numpy==1.24.3  
pandas==2.2.3  
matplotlib==3.10.1  
scikit-learn==1.6.1  
joblib==1.5.0
```

Instructions to Run the Test Script:

1. Set up the Python environment. See requirements.txt for dependencies.
2. Run the test script for Model 1 using:
`python proj3_classification_test.py --model <model_path> --test_csv <test_data_csv_path>`

eg.

```
python proj3_classification_test.py --model classification.keras --test_csv mushrooms_test.csv
```

3. Run the test script for Model 2 using:

```
python proj3_extractSVM_test.py --model <model_path> --test_csv <test_data_csv_path>
--feature_model <feature_model_path>
```

eg.

```
python proj3_extractSVM_test.py --model svm_classifier.pkl --feature_model
feature_extractor.keras
```

3. Output - Model accuracy

Note: Ensure the model_path and test_data_csv_path are correct before running the script.